

Shiva Preetham Chinthalachervu

📞 612-801-7170 — ✉ shivacheruvu6@gmail.com — 🌐 github.com/shivacheruvu

Summary — Data Engineer experienced in building scalable PySpark pipelines, modernizing legacy enterprise data warehouses, and optimizing cloud-native architecture on Databricks, AWS, and GCP. Skilled in forward-deployed environments, directly collaborating with business stakeholders to resolve complex data anomalies, establish automated validation frameworks, and engineer query-ready Delta Lake models that accelerate business intelligence and analytics.

Skills

Languages Python, PySpark, SQL, Bash, Java
Data Eng Databricks, Delta Lake, ETL/ELT, Spark
Cloud AWS (S3, EC2), GCP (GCS), Azure

Databases Oracle, Informix, PostgreSQL, Snowflake
Orchestration Databricks Workflows, YAML, CI/CD
Tools Git, Antigravity IDE, Cursor, OneTrust

Experience

Frontier/Verizon

Mar 2025 – Present

Data Engineer (Forward Deployed)

- Architected and deployed scalable PySpark pipelines on Databricks to replace legacy Informatica (IICS) and Oracle ETL workflows, cutting pipeline runtime by 42% and successfully decommissioning legacy infrastructure
- Embedded directly with telecommunications stakeholders to engineer custom Databricks POCs, translating complex operational requirements into automated tracking frameworks and production-ready data pipelines
- Designed store-and-forward batch ingestion frameworks processing multi-gigabyte compressed .gz files into Delta Lake medallion architectures, optimizing partition pruning and Z-Ordering to accelerate queries by 3.5x
- Built an automated production diff and reconciliation engine using PySpark to perform primary key validation, schema drift detection, and discrepancy logging across billions of records between on-prem EDW and cloud storage
- Configured modular Databricks Multi-Task Workflows utilizing YAML-based task definitions, conditional job execution, and automated alerting protocols to achieve 99.9% pipeline SLA reliability
- Implemented end-to-end OneTrust and CCPA automated deletion pipelines across Databricks Unity Catalog storage volumes and object stores, ensuring 100% enterprise regulatory compliance

The Insight INC

Dec 2024 – Mar 2025

Technical Data Analyst

- Led cross-system data reconciliation, writing complex SQL and PySpark scripts to investigate log file count mismatches, identify data type conversion errors, and resolve severe timestamp filtering anomalies
- Engineered highly optimized, query-ready data models specifically structured for downstream dashboard exploration, directly accelerating business intelligence query response times for executive stakeholders
- Developed a dynamic tracking framework that ingested unstructured operational logs into consolidated reporting tables, generating automated KPI discrepancy reports to monitor end-to-end pipeline health
- Wrote advanced SQL queries utilizing complex date-filtering options and window functions to isolate missing historical records, executing targeted backfill procedures to ensure 100% financial data completeness

Education

University of Minnesota

Bachelor of Science in Data Science

Certifications

- **Databricks Certified Professional Data Engineer** – (In Progress)
- **Google Cloud Certified Professional Data Engineer** – (In Progress)
- **AWS Certified Cloud Practitioner** – Foundational expertise in cloud architecture and secure deployment
- **Okta Certified Professional** – Competency in identity access management and application security

Projects

Automated Schema & Primary Key Extraction Utility

Feb 2024 – Mar 2024

- Developed a dynamic PySpark utility for programmatic schema column searching, metadata extraction, and primary key identification, eliminating hardcoded dependencies and cutting team debugging time by over 50%
- Integrated extraction scripts into a CI/CD pipeline, configuring YAML workflows to send automated, formatted notification emails containing schema drift alerts directly to on-call engineering teams